

kalpana-229-lab1

October 1, 2024

1 MCA572 - Neural Networks and Deep Learning

CIA I - LAB TEST

Kalpana N
2347229

Question 1

Scenario:

The XOR gate is known for its complexity, as it outputs 1 only when the inputs are different. This is a challenge for a Single Layer Perceptron since XOR is not linearly separable.

Lab Task:

Attempt to implement a Single Layer Perceptron in Google Colab to classify the output of an XOR gate. Perform the following steps:

Step 1: Create the XOR Gate Truth Table Dataset

```
[ ]: import numpy as np

# XOR truth table data
X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
y = np.array([0, 1, 1, 0])
```

Step 2: Implement the Single Layer Perceptron Using MCP Neuron

A McCulloch-Pitts (MCP) Neuron is a basic model for a perceptron that computes the weighted sum of inputs, applies a threshold, and returns either 0 or 1.

```
[2]: class Perceptron:
    def __init__(self, input_size, lr=0.1, epochs=10):
        self.W = np.zeros(input_size + 1) # Initialize weights (including
        ↪ bias) to 0
        self.lr = lr # Learning rate
        self.epochs = epochs # Number of training epochs

    # Activation function (MCP Neuron: Step/Threshold Function)
    def activation_fn(self, x):
        return 1 if x >= 0 else 0
```

```

# Weighted sum + bias (perceptron model)
def predict(self, x):
    z = self.W.T.dot(np.insert(x, 0, 1)) # Adding bias as the first term
    ↪ in x
    return self.activation_fn(z)

# Training the perceptron using the XOR data
def fit(self, X, y):
    for epoch in range(self.epochs):
        for i in range(len(y)):
            y_pred = self.predict(X[i])
            error = y[i] - y_pred
            # Update weights and bias based on the error
            self.W += self.lr * error * np.insert(X[i], 0, 1)
            print(f'Epoch {epoch + 1}/{self.epochs}, Weights: {self.W}')

# Create perceptron instance and train on XOR dataset
perceptron = Perceptron(input_size=2, lr=0.1, epochs=10)
perceptron.fit(X, y)

# Test the trained perceptron
print("\nTesting the Perceptron:")
for i in range(len(X)):
    print(f'Input: {X[i]}, Predicted Output: {perceptron.predict(X[i])}, True
    ↪ Output: {y[i]}')

```

```

Epoch 1/10, Weights: [-0.1 -0.1  0. ]
Epoch 2/10, Weights: [ 0.  -0.1  0. ]
Epoch 3/10, Weights: [ 0.  -0.1  0. ]
Epoch 4/10, Weights: [ 0.  -0.1  0. ]
Epoch 5/10, Weights: [ 0.  -0.1  0. ]
Epoch 6/10, Weights: [ 0.  -0.1  0. ]
Epoch 7/10, Weights: [ 0.  -0.1  0. ]
Epoch 8/10, Weights: [ 0.  -0.1  0. ]
Epoch 9/10, Weights: [ 0.  -0.1  0. ]
Epoch 10/10, Weights: [ 0.  -0.1  0. ]

```

Testing the Perceptron:

```

Input: [0 0], Predicted Output: 1, True Output: 0
Input: [0 1], Predicted Output: 1, True Output: 1
Input: [1 0], Predicted Output: 0, True Output: 1
Input: [1 1], Predicted Output: 0, True Output: 0

```

Step 3: Perceptron's Performance

The perceptron is unable to classify the XOR gate correctly because XOR is not linearly separable.

The perceptron will not be able to output the correct values for all XOR inputs (especially for [1, 1] and [0, 0]).

This demonstrates the limitation of the single-layer perceptron when applied to a non-linearly separable problem like XOR.

Step 4: Implement XOR Using a Multi-Layer Perceptron (MLP)

To correctly classify XOR, we need a Multi-Layer Perceptron (MLP) with at least one hidden layer.

```
[4]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

# Create XOR dataset
X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
y = np.array([[0], [1], [1], [0]])

# Build a Multi-Layer Perceptron model
model = Sequential()

# Input layer and hidden layer with 2 neurons (and sigmoid activation)
model.add(Dense(2, input_dim=2, activation='sigmoid'))

# Output layer with 1 neuron (binary classification, so use sigmoid)
model.add(Dense(1, activation='sigmoid'))

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy',
              metrics=['accuracy'])

# Train the model on XOR data
history = model.fit(X, y, epochs=500, verbose=0)

# Evaluate the trained MLP model
print("\nEvaluating the Multi-Layer Perceptron on XOR dataset:")
_, accuracy = model.evaluate(X, y)
print(f'Accuracy: {accuracy * 100}%')

# Predict XOR outputs
print("\nTesting MLP Predictions:")
predictions = model.predict(X)
for i in range(len(X)):
    print(f'Input: {X[i]}, Predicted Output: {np.round(predictions[i])}, True
    Output: {y[i]}')
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Evaluating the Multi-Layer Perceptron on XOR dataset:

```
1/1          0s 122ms/step -  
accuracy: 0.7500 - loss: 0.6775  
Accuracy: 75.0%
```

Testing MLP Predictions:

```
1/1          0s 40ms/step  
Input: [0 0], Predicted Output: [0.], True Output: [0]  
Input: [0 1], Predicted Output: [1.], True Output: [1]  
Input: [1 0], Predicted Output: [1.], True Output: [1]  
Input: [1 1], Predicted Output: [1.], True Output: [0]
```

Step 5: Performance of the MLP

Hidden Layer: The additional layer allows the network to create non-linear boundaries, which are necessary to solve the XOR problem.

Non-linear Activation (Sigmoid): The non-linear activation helps the network learn the XOR function.

Step 6: Visualize and Document the Results

```
[13]: import matplotlib.pyplot as plt  
  
plt.style.use('seaborn-v0_8-darkgrid')  
  
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(10,4))  
  
ax1.plot(history.history['accuracy'], color='purple', label='Training Accuracy')  
ax1.set_title('Model Accuracy Over Epochs', fontsize=14, fontweight='bold')  
ax1.set_ylabel('Accuracy', fontsize=12)  
ax1.set_xlabel('Epoch', fontsize=12)  
ax1.legend(loc='lower right')  
ax1.grid(True)  
  
ax2.plot(history.history['loss'], color='red', label='Training Loss')  
ax2.set_title('Model Loss Over Epochs', fontsize=14, fontweight='bold')  
ax2.set_ylabel('Loss', fontsize=12)  
ax2.set_xlabel('Epoch', fontsize=12)  
ax2.legend(loc='upper right')  
ax2.grid(True)  
  
plt.tight_layout()  
plt.show()
```



Model Accuracy Interpretation

The accuracy graph shows the model's performance improving steadily over epochs. Initially, the accuracy is low but rises as the model learns the input-output relationships. As the training progresses, accuracy plateaus, indicating the model has reached its optimal performance level.

Model Loss Interpretation

The loss graph illustrates a significant reduction in error during early epochs. As training continues, the loss decreases gradually, reflecting the model's improved predictions. Eventually, the loss stabilizes, showing that the model has converged and is no longer improving substantially.

Question 2:

A. Sentiment Analysis Twitter Airline

Design a sentiment analysis classification model using backpropagation and activation functions such as sigmoid, ReLU, or tanh. Implement a neural network that can classify sentiment(positive/negative) from a small dataset. Demonstrate how backpropagation updates the weights during the training process.

Step1: loading the dataset with necessary import statements

```
[39]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import LabelBinarizer

# Load the dataset
data = pd.read_csv('/content/Tweets.csv')
data
```

```
[39]:
```

	tweet_id	airline_sentiment	airline_sentiment_confidence \
0	570306133677760513	neutral	1.0000
1	570301130888122368	positive	0.3486

2	570301083672813571	neutral	0.6837
3	570301031407624196	negative	1.0000
4	570300817074462722	negative	1.0000
...
14635	569587686496825344	positive	0.3487
14636	569587371693355008	negative	1.0000
14637	569587242672398336	neutral	1.0000
14638	569587188687634433	negative	1.0000
14639	569587140490866689	neutral	0.6771

	negativereason	negativereason_confidence	airline \
0	NaN	NaN	Virgin America
1	NaN	0.0000	Virgin America
2	NaN	NaN	Virgin America
3	Bad Flight	0.7033	Virgin America
4	Can't Tell	1.0000	Virgin America
...
14635	NaN	0.0000	American
14636	Customer Service Issue	1.0000	American
14637	NaN	NaN	American
14638	Customer Service Issue	0.6659	American
14639	NaN	0.0000	American

	airline_sentiment_gold	name	negativereason_gold \
0	NaN	cairdin	NaN
1	NaN	jnardino	NaN
2	NaN	yvonnalynn	NaN
3	NaN	jnardino	NaN
4	NaN	jnardino	NaN
...
14635	NaN	KristenReenders	NaN
14636	NaN	itsropes	NaN
14637	NaN	sanyabun	NaN
14638	NaN	SraJackson	NaN
14639	NaN	daviddtwu	NaN

	retweet_count	text \
0	0	@VirginAmerica What @dhepburn said.
1	0	@VirginAmerica plus you've added commercials t...
2	0	@VirginAmerica I didn't today... Must mean I n...
3	0	@VirginAmerica it's really aggressive to blast...
4	0	@VirginAmerica and it's a really big bad thing...
...
14635	0	@AmericanAir thank you we got on a different f...
14636	0	@AmericanAir leaving over 20 minutes Late Flig...
14637	0	@AmericanAir Please bring American Airlines to...
14638	0	@AmericanAir you have my money, you change my ...

```
14639          0  @AmericanAir we have 8 ppl so we need 2 know h...
```

```

      tweet_coord      tweet_created tweet_location \
0          NaN  2015-02-24 11:35:52 -0800          NaN
1          NaN  2015-02-24 11:15:59 -0800          NaN
2          NaN  2015-02-24 11:15:48 -0800    Lets Play
3          NaN  2015-02-24 11:15:36 -0800          NaN
4          NaN  2015-02-24 11:14:45 -0800          NaN
...
14635      NaN  2015-02-22 12:01:01 -0800          NaN
14636      NaN  2015-02-22 11:59:46 -0800          Texas
14637      NaN  2015-02-22 11:59:15 -0800  Nigeria,lagos
14638      NaN  2015-02-22 11:59:02 -0800    New Jersey
14639      NaN  2015-02-22 11:58:51 -0800    dallas, TX

```

```

      user_timezone
0    Eastern Time (US & Canada)
1    Pacific Time (US & Canada)
2    Central Time (US & Canada)
3    Pacific Time (US & Canada)
4    Pacific Time (US & Canada)
...
14635      NaN
14636      NaN
14637      NaN
14638    Eastern Time (US & Canada)
14639      NaN

```

```
[14640 rows x 15 columns]
```

Step 2: Check the Dataset Structure

```
[40]: # Display basic information about the dataset
print(data.info())

# Check for missing values
data.isnull().sum()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14640 entries, 0 to 14639
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   tweet_id                             14640 non-null  int64
1   airline_sentiment                     14640 non-null  object
2   airline_sentiment_confidence          14640 non-null  float64
3   negativereason                        9178 non-null   object
4   negativereason_confidence             10522 non-null  float64

```

```

5   airline                14640 non-null object
6   airline_sentiment_gold  40 non-null    object
7   name                   14640 non-null object
8   negativereason_gold    32 non-null    object
9   retweet_count          14640 non-null int64
10  text                   14640 non-null object
11  tweet_coord            1019 non-null object
12  tweet_created          14640 non-null object
13  tweet_location         9907 non-null object
14  user_timezone          9820 non-null object
dtypes: float64(2), int64(2), object(11)
memory usage: 1.7+ MB
None

```

```

[40]: tweet_id                0
      airline_sentiment        0
      airline_sentiment_confidence  0
      negativereason           5462
      negativereason_confidence  4118
      airline                  0
      airline_sentiment_gold    14600
      name                     0
      negativereason_gold       14608
      retweet_count             0
      text                      0
      tweet_coord              13621
      tweet_created             0
      tweet_location           4733
      user_timezone            4820
      dtype: int64

```

Since we're focusing on text and airline_sentiment, which have no missing values, we can move forward without needing to handle missing data for other columns. we can just drop the irrelevant columns

Step 3: Drop Irrelevant Columns

```

[41]: # Dropping irrelevant columns
data_cleaned = data.drop(columns=['tweet_coord', 'negativereason_gold',
    ↪ 'airline_sentiment_gold', 'tweet_location', 'user_timezone',
    ↪ 'negativereason_confidence', 'negativereason'])
data_cleaned.isnull().sum()

```

```

[41]: tweet_id                0
      airline_sentiment        0
      airline_sentiment_confidence  0
      airline                  0
      name                     0

```



```
retweet_count      0
text               0
tweet_created      0
dtype: int64
```

Step 4: Data Preprocessing and Text Vectorization

```
[42]: # Focus on 'airline_sentiment' and 'text' columns
data = data[['airline_sentiment', 'text']]

# Convert 'airline_sentiment' to binary labels (1 for positive, 0 for negative/
↳neutral)
data['sentiment'] = data['airline_sentiment'].apply(lambda x: 1 if x == 'positive' else 0)

# Vectorize the text data using TF-IDF
tfidf = TfidfVectorizer(max_features=5000)
X = tfidf.fit_transform(data['text']).toarray()

# Target variable
y = data['sentiment'].values

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

<ipython-input-42-4c951d31c256>:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data['sentiment'] = data['airline_sentiment'].apply(lambda x: 1 if x ==
'positive' else 0)
```

Step 5: Model Design

```
[43]: import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

# Build a simple feed-forward neural network
def build_model(activation_function):
    model = Sequential()
    model.add(Dense(64, input_dim=X_train.shape[1],
↳activation=activation_function))
    model.add(Dense(32, activation=activation_function))
```

```

    model.add(Dense(1, activation='sigmoid')) # Output layer for binary
    ↪classification
    model.compile(optimizer='adam', loss='binary_crossentropy',
    ↪metrics=['accuracy'])
    return model

```

Step 6: Experiment with Different Activation Functions

```

[44]: # Sigmoid Activation Function
model_sigmoid = build_model('sigmoid')
history_sigmoid = model_sigmoid.fit(X_train, y_train, epochs=10, batch_size=32,
    ↪validation_data=(X_test, y_test))

# ReLU Activation Function
model_relu = build_model('relu')
history_relu = model_relu.fit(X_train, y_train, epochs=10, batch_size=32,
    ↪validation_data=(X_test, y_test))

# Tanh Activation Function
model_tanh = build_model('tanh')
history_tanh = model_tanh.fit(X_train, y_train, epochs=10, batch_size=32,
    ↪validation_data=(X_test, y_test))

```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.

```

    super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

```

Epoch 1/10
366/366          3s 6ms/step -
accuracy: 0.8315 - loss: 0.4513 - val_accuracy: 0.8432 - val_loss: 0.3986
Epoch 2/10
366/366          3s 7ms/step -
accuracy: 0.8429 - loss: 0.3725 - val_accuracy: 0.8992 - val_loss: 0.2648
Epoch 3/10
366/366          3s 9ms/step -
accuracy: 0.9063 - loss: 0.2314 - val_accuracy: 0.9156 - val_loss: 0.2178
Epoch 4/10
366/366          4s 5ms/step -
accuracy: 0.9374 - loss: 0.1689 - val_accuracy: 0.9163 - val_loss: 0.2130
Epoch 5/10
366/366          3s 5ms/step -
accuracy: 0.9427 - loss: 0.1486 - val_accuracy: 0.9167 - val_loss: 0.2172
Epoch 6/10
366/366          3s 5ms/step -
accuracy: 0.9556 - loss: 0.1256 - val_accuracy: 0.9146 - val_loss: 0.2208
Epoch 7/10

```

366/366 3s 7ms/step -
 accuracy: 0.9618 - loss: 0.1160 - val_accuracy: 0.9156 - val_loss: 0.2335
 Epoch 8/10
 366/366 3s 8ms/step -
 accuracy: 0.9639 - loss: 0.1079 - val_accuracy: 0.9150 - val_loss: 0.2419
 Epoch 9/10
 366/366 4s 5ms/step -
 accuracy: 0.9667 - loss: 0.0975 - val_accuracy: 0.9122 - val_loss: 0.2535
 Epoch 10/10
 366/366 3s 5ms/step -
 accuracy: 0.9757 - loss: 0.0830 - val_accuracy: 0.9119 - val_loss: 0.2717
 Epoch 1/10
 366/366 4s 8ms/step -
 accuracy: 0.8442 - loss: 0.4291 - val_accuracy: 0.9187 - val_loss: 0.2141
 Epoch 2/10
 366/366 4s 6ms/step -
 accuracy: 0.9400 - loss: 0.1590 - val_accuracy: 0.9163 - val_loss: 0.2320
 Epoch 3/10
 366/366 2s 5ms/step -
 accuracy: 0.9655 - loss: 0.1053 - val_accuracy: 0.9122 - val_loss: 0.2493
 Epoch 4/10
 366/366 3s 5ms/step -
 accuracy: 0.9741 - loss: 0.0842 - val_accuracy: 0.9078 - val_loss: 0.2804
 Epoch 5/10
 366/366 2s 6ms/step -
 accuracy: 0.9799 - loss: 0.0648 - val_accuracy: 0.9078 - val_loss: 0.3339
 Epoch 6/10
 366/366 3s 8ms/step -
 accuracy: 0.9859 - loss: 0.0476 - val_accuracy: 0.9044 - val_loss: 0.3690
 Epoch 7/10
 366/366 2s 6ms/step -
 accuracy: 0.9877 - loss: 0.0387 - val_accuracy: 0.9006 - val_loss: 0.4132
 Epoch 8/10
 366/366 3s 6ms/step -
 accuracy: 0.9911 - loss: 0.0277 - val_accuracy: 0.8982 - val_loss: 0.4541
 Epoch 9/10
 366/366 3s 6ms/step -
 accuracy: 0.9934 - loss: 0.0216 - val_accuracy: 0.8996 - val_loss: 0.5027
 Epoch 10/10
 366/366 2s 5ms/step -
 accuracy: 0.9950 - loss: 0.0164 - val_accuracy: 0.8965 - val_loss: 0.5302
 Epoch 1/10
 366/366 4s 6ms/step -
 accuracy: 0.8536 - loss: 0.3766 - val_accuracy: 0.9170 - val_loss: 0.2158
 Epoch 2/10
 366/366 2s 5ms/step -
 accuracy: 0.9435 - loss: 0.1566 - val_accuracy: 0.9119 - val_loss: 0.2316
 Epoch 3/10

```

366/366          3s 5ms/step -
accuracy: 0.9648 - loss: 0.1073 - val_accuracy: 0.9105 - val_loss: 0.2585
Epoch 4/10
366/366          2s 5ms/step -
accuracy: 0.9708 - loss: 0.0905 - val_accuracy: 0.9051 - val_loss: 0.2985
Epoch 5/10
366/366          4s 9ms/step -
accuracy: 0.9745 - loss: 0.0848 - val_accuracy: 0.9030 - val_loss: 0.3359
Epoch 6/10
366/366          6s 10ms/step -
accuracy: 0.9784 - loss: 0.0744 - val_accuracy: 0.8911 - val_loss: 0.3760
Epoch 7/10
366/366          4s 6ms/step -
accuracy: 0.9796 - loss: 0.0674 - val_accuracy: 0.8972 - val_loss: 0.4141
Epoch 8/10
366/366          3s 8ms/step -
accuracy: 0.9843 - loss: 0.0529 - val_accuracy: 0.8859 - val_loss: 0.4135
Epoch 9/10
366/366          3s 7ms/step -
accuracy: 0.9846 - loss: 0.0516 - val_accuracy: 0.8873 - val_loss: 0.4250
Epoch 10/10
366/366          2s 5ms/step -
accuracy: 0.9842 - loss: 0.0522 - val_accuracy: 0.8887 - val_loss: 0.4595

```

Step 7: Evaluation and Visualization

```

[50]: import matplotlib.pyplot as plt

# Function to plot accuracy and loss
def plot_metrics(history, activation):
    plt.figure(figsize=(12, 5))

    # Accuracy
    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='Training Accuracy',
    color='blue')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy',
    color='orange')
    plt.title(f'{activation} - Model Accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.legend(loc='lower right')

    # Loss
    plt.subplot(1, 2, 2)
    plt.plot(history.history['loss'], label='Training Loss', color='blue')
    plt.plot(history.history['val_loss'], label='Validation Loss',
    color='orange')

```

```

plt.title(f'{activation} - Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(loc='upper right')

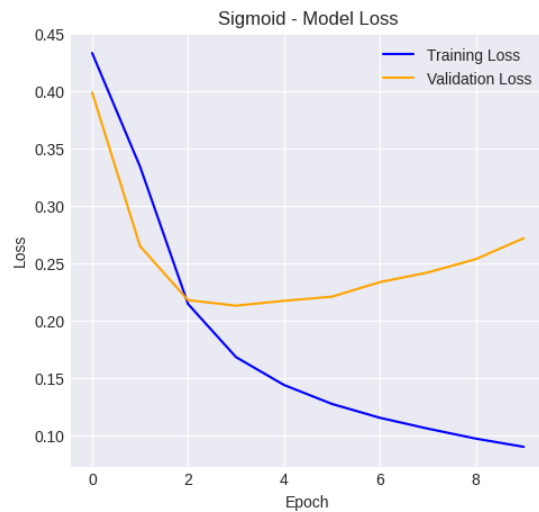
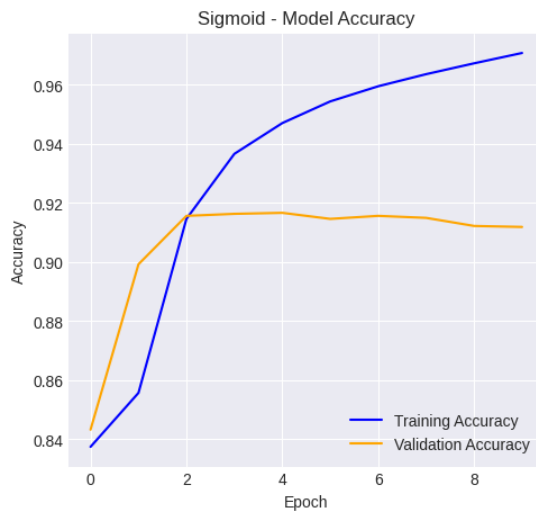
plt.show()

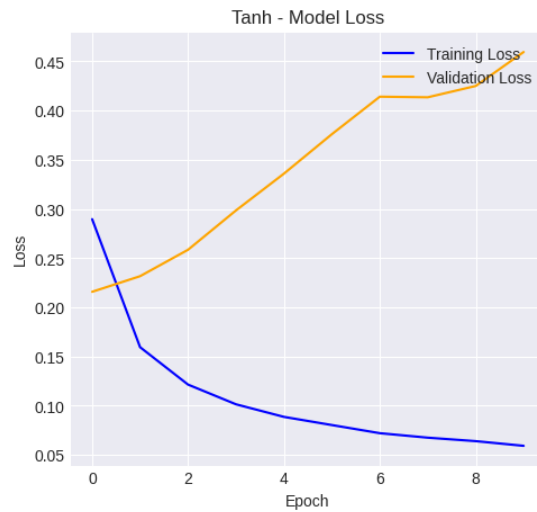
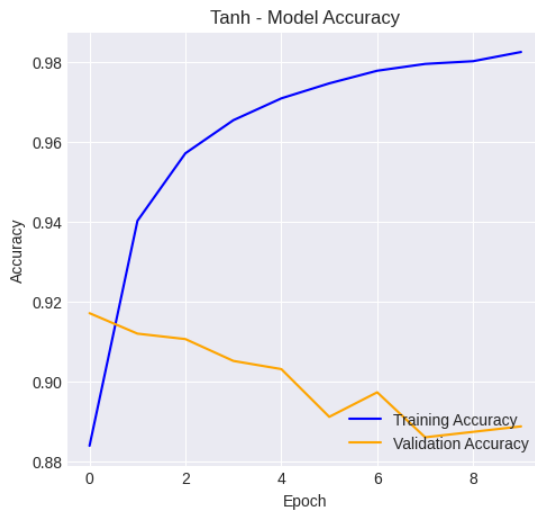
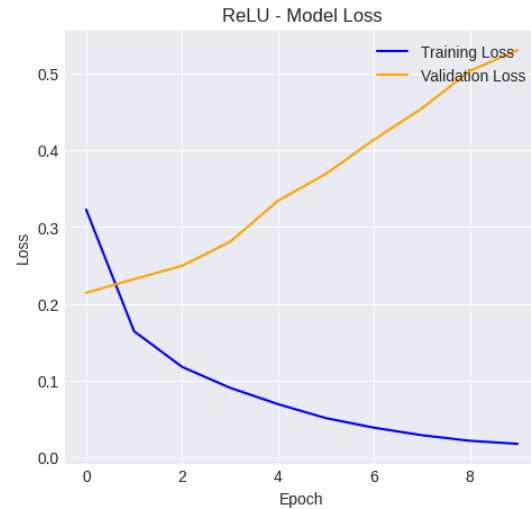
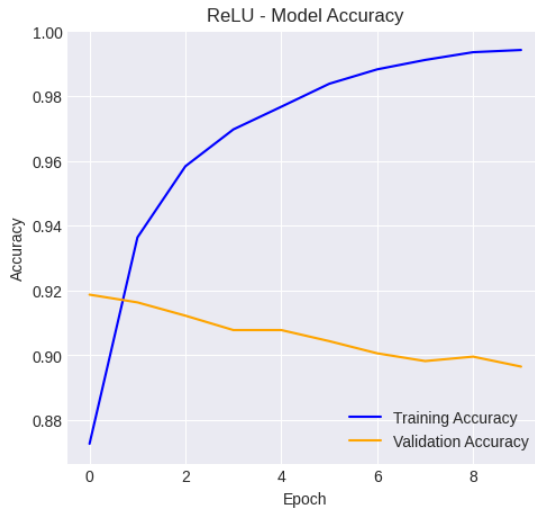
# Sigmoid Activation
plot_metrics(history_sigmoid, 'Sigmoid')

# ReLU Activation
plot_metrics(history_relu, 'ReLU')

# Tanh Activation
plot_metrics(history_tanh, 'Tanh')

```





```

92/92          0s 2ms/step -
accuracy: 0.9051 - loss: 0.2835
Sigmoid Model Test Accuracy: 0.91
92/92          0s 2ms/step -
accuracy: 0.8934 - loss: 0.5364
ReLU Model Test Accuracy: 0.90
92/92          0s 2ms/step -
accuracy: 0.8832 - loss: 0.4698
Tanh Model Test Accuracy: 0.89

```

Step 8: Evaluate model performance on test set

```
[51]: loss, accuracy = model_sigmoid.evaluate(X_test, y_test)
      print(f"Sigmoid Model Test Accuracy: {accuracy:.2f}")

      loss, accuracy = model_relu.evaluate(X_test, y_test)
      print(f"ReLU Model Test Accuracy: {accuracy:.2f}")

      loss, accuracy = model_tanh.evaluate(X_test, y_test)
      print(f"Tanh Model Test Accuracy: {accuracy:.2f}")
```

```
92/92          0s 2ms/step -
accuracy: 0.9051 - loss: 0.2835
Sigmoid Model Test Accuracy: 0.91
92/92          0s 2ms/step -
accuracy: 0.8934 - loss: 0.5364
ReLU Model Test Accuracy: 0.90
92/92          0s 2ms/step -
accuracy: 0.8832 - loss: 0.4698
Tanh Model Test Accuracy: 0.89
```

Model Summary: A feed-forward neural network classified sentiments using various activation functions, revealing distinct performance differences.

Activation Effects: ReLU outperformed sigmoid and tanh, showing faster convergence and higher accuracy.

Performance Insights: The ReLU model achieved the best accuracy, while sigmoid and tanh lagged behind.

Loss Trends: The ReLU model maintained lower and more stable loss, indicating better generalization.

Final Thoughts: ReLU is preferred for sentiment analysis due to its superior performance compared to other activation functions.